

# The Relational Model

## ER to Relational / SQL

# Announcements

- Waiting list: all done (sorry)
- Project 1: Due Thursday morning
  - Late submissions: Under Eugene Wu's door (421 Mudd in the DSI space)

# Review

Relation: Set of tuples with typed values (table)

Schema: Names and types for values in relation

Database: Set of relations

SQL: Structured Query Language

Integrity constraint: Restrictions on valid data

Candidate key: Minimal fields that identify tuple

Primary key: Designated identifier for a tuple

Foreign key: Reference to another tuple

# Referential Integrity

A database instance has *referential integrity* if all foreign key constraints are enforced no dangling references

Examples where referential integrity is not enforced

- HTML links

- Yellow page listing

- Restaurant menus

- Some relational databases!

# How to Enforce Integrity Constraints

Run checks anytime database changes

On INSERT

what if new Enrolled tuple refers to non-existent student?

Reject insertion

On DELETE (many options)

what if Students tuple is deleted?

delete dependent Enrolled tuples

reject deletion

set Enrolled.sid to default value or null

(null means 'unknown' or 'inapplicable' in SQL)

# General Constraints

Boolean constraint expression added to schema

Very powerful; not widely used

```
CREATE TABLE Enrolled(  
    sid int,  
    cid int,  
    grade char(2),  
    CHECK (  
        grade = 'A' or grade = 'B' or  
        grade = 'C' or grade = 'D' or  
        grade = 'F'  
    )  
)
```

# Where do ICs come from?


Based on application semantics and use cases

IC is statement about the world (all possible instances)

Can't infer ICs by staring at an instance

e.g. Is “login” a unique candidate key? Maybe?

Unique and foreign key ICs are very common  
(others less so)




Translating ER → Relational Models



# Entity Set → Relation

Include all attributes

Entity set key become relation primary key



```
CREATE TABLE Course(  
  cid int,  
  name text,  
  loc text,  
  PRIMARY KEY (cid)  
)
```


# Relationship Set w/out constraint → Relation

Relation includes:

All attributes of the relationship set

Add keys for each entity set as foreign keys: *superkey*

(if there are other constraints: may not need all the keys)



```
CREATE TABLE Takes(  
  uid int,  
  cid int,  
  since date,  
  PRIMARY KEY (uid, cid),  
  FOREIGN KEY (uid) REFERENCES Users,  
  FOREIGN KEY (cid) REFERENCES Courses  
)
```


# At Most One → Relation

Add relationship attributes (none here)

Add keys for entity set as foreign keys

What is the primary key?

Only cid: Zero or 1 Course in this relationship



```
CREATE TABLE Instructs(  
    uid int,  
    cid int,  
    PRIMARY KEY (cid),  
    FOREIGN KEY (uid) REFERENCES Users,  
    FOREIGN KEY (cid) REFERENCES Courses  
);
```


# At Most One: Combine?

Zero or 1 courses, cid is primary key; Similar to ???

Combine Instructs attributes into Courses (preferred)

How to represent courses without instructor?


NULL value for uid (and other Instructs attributes)



```
CREATE TABLE Course_Instructs(  
  cid int  
  uid int, Rename instructor_uid?  
  name text,  
  loc text,  
  PRIMARY KEY (cid),  
  FOREIGN KEY (uid) REFERENCES Users  
)
```

# Key Constraint → Relation

- How to translate this ER diagram?
- NOTE: I'm pretty sure this is wrong! Hence hidden



```
CREATE TABLE Course_Instructs(  
    cid int  
    uid int,  
    name text,  
    loc text,  
    username text,  
    age text,  
    PRIMARY KEY (cid),  
    UNIQUE (uid)  
)
```


# Exactly One Constraint → Relation

Represent Course must have Instructor?

Combine relationship into Courses; NOT NULL

What happens if we delete User who is Instructor?


Default ON DELETE NO ACTION: Don't permit



```
CREATE TABLE Course_Instructs(  
  cid int  
  uid int, NOT NULL,  
  name text,  
  loc text,  
  PRIMARY KEY (cid),  
  FOREIGN KEY (uid) REFERENCES Users  
)  
  ON DELETE NO ACTION  
)
```

# At Least One Constraint?

```
CREATE TABLE Instructs(  
  cid int,  
  uid int NOT NULL,  
  PRIMARY KEY (cid, uid),  
  FOREIGN KEY (cid) REFERENCES Course,  
  FOREIGN KEY (uid) REFERENCES User)
```




Same as no constraint!  
Can't be easily expressed

# Weak Entity → Relation

Weak entity set and identifying relationship set are translated into a single table.

When the owner entity is deleted, all owned weak entities must also be deleted.



```
CREATE TABLE Wall_Posted(  
    pid int  
    post_text text,  
    posted_time DATE,  
    uid int,  
    PRIMARY KEY (pid, uid),  
    FOREIGN KEY (uid) REFERENCES Users  
    ON DELETE CASCADE  
)
```



# ISA Hierarchies

## Option 1: Keep base relation

Instructors & Students recorded in Users

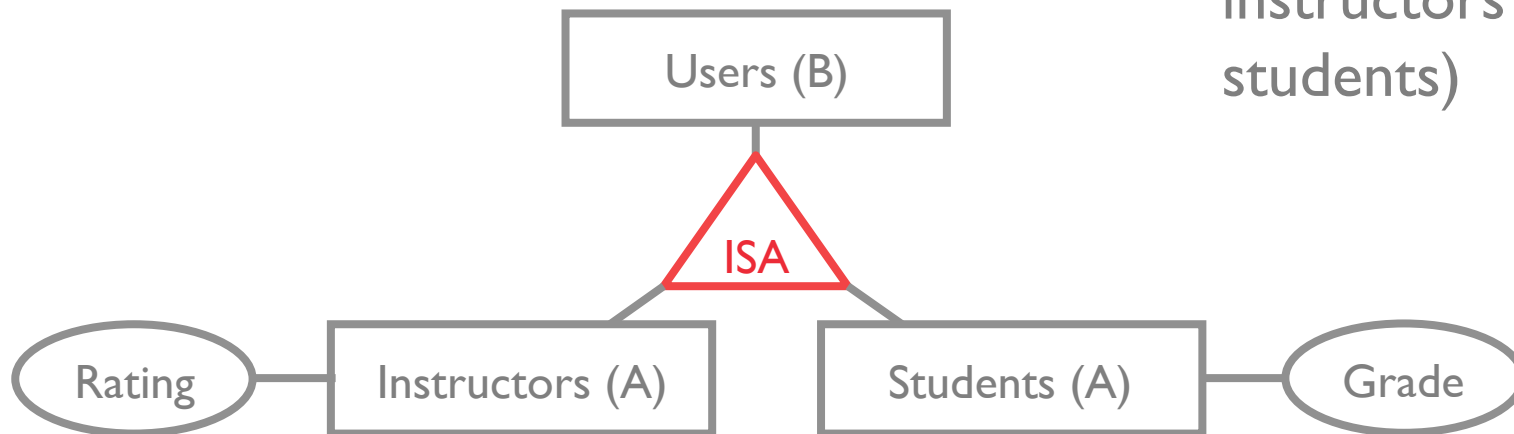
Extra info in Instructors or Students relation

Primary key = key of parent table; is foreign key for child tables

## Option 2: Only keep child relations

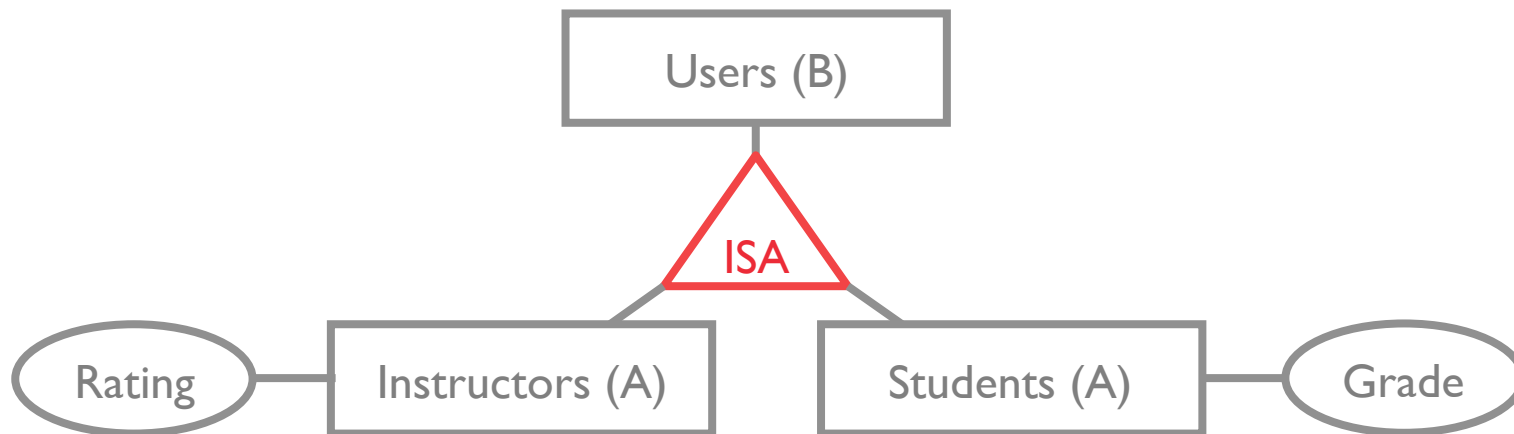
Children copy attributes from Users

Only if covering  
constraint = yes  
(all users are  
instructors or  
students)

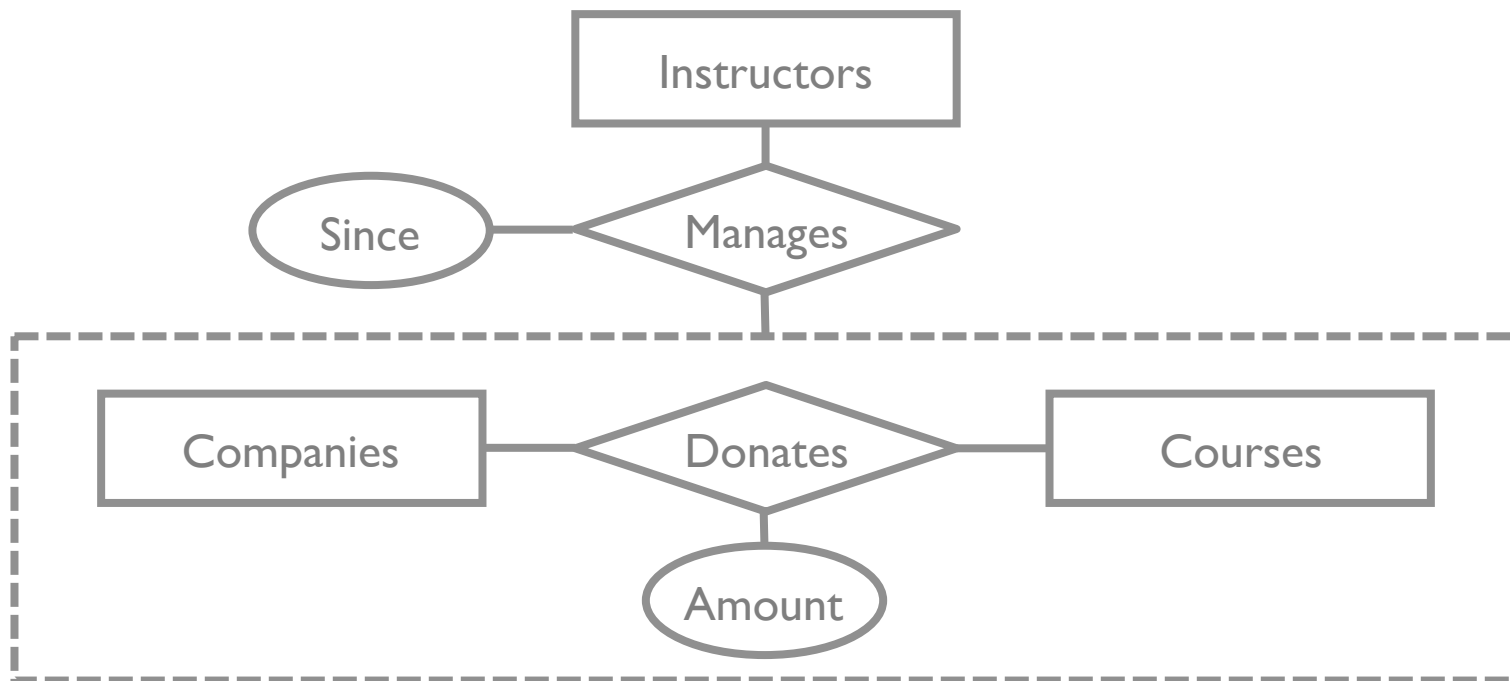


# ISA Hierarchies

```
CREATE TABLE Users(uid int, name text, PRIMARY KEY(uid))  
CREATE TABLE Instructors(uid int, rating int,  
    PRIMARY KEY(uid), FOREIGN KEY (uid) REFERENCES Users)  
CREATE TABLE Students(uid int, grade char(2),  
    PRIMARY KEY(uid), FOREIGN KEY (uid) REFERENCES Users)
```




# Aggregation



# Aggregation

Convert the aggregated relationship into an entity  
Convert like any other entity

E.g. Donates: PRIMARY KEY (Company\_id, Course\_id)  
Manages: References this key



# DML: Introduction to Queries

Query: Ask a question

Declarative: ask what you want, now how

Queries are high level, “readable”

DBMS decides how to make it fast

Precise semantics for relational queries

Lets DBMS choose different ways to run query while ensuring answer is the same

# INSERT/DELETE

## Add a tuple

```
INSERT INTO Students(sid, name, login, age, gpa)
VALUES (4, 'wu', 'wu@cs', 20, 5)
```

## Delete tuples satisfying a predicate (condition)

```
DELETE FROM Students S
WHERE S.name = 'wu'
```

# Basic Single Table SELECT

SELECT (*output*) FROM (*input*)  
WHERE (*condition*)

SELECT \* FROM Students

SELECT name FROM Students

SELECT name FROM Students WHERE age < 21

SELECT name, login FROM Students WHERE gpa >= 3

sid	name	login	age	gpa
1	eugene	ewu@cs	20	2.5
2	luis	luis@cs	20	3.5
3	ken	ken@math	33	3.9

# Ambiguous names

E.g. Students: (sid, name, ...)

Enrolled: (sid, cid, grade)

*Qualified* names: Use table name: Students.age

Rename: AS (optional): shortcuts, ambiguity, clarity

```
SELECT S.sid, S.name FROM Students AS S
```

```
SELECT S.sid, S.name FROM Students S
```



# Single Table Semantics

- A *conceptual evaluation method* for previous query:
  - 1. FROM clause: retrieve Students relation
  - 2. WHERE clause: Check conditions, discard tuples that fail
  - 3. SELECT clause: Delete unwanted fields
- Remember, this is *conceptual*. Actual evaluation will be *much* more efficient, but must produce the same answers.

# Related data: Multiple tables

What does this return?

```
SELECT S.name, E.cid  
FROM Students S, Enrolled E  
WHERE S.sid = E.sid AND  
       E.grade = 'A'
```

Enrolled

sid	cid	grade
1	2	A
1	3	B
2	2	A+

Students

sid	name
1	eugene
2	luis

Result

name	cid
eugene	2

# Multi-Table Semantics

- Modify the FROM clause evaluation
  - 1. FROM clause: compute *cross-product* of Students and

Enrolled

sid	cid	grade
1	2	A
1	3	B
2	2	A+

Students

sid	name
1	eugene
2	luis

Cross-product

sid	cid	grade	sid	name
1	2	A	1	eugene
1	3	B	1	eugene
2	2	A+	1	eugene
1	2	A	2	luis
1	3	B	2	luis
2	2	A+	2	luis

# Multi-Table Semantics

- Modify the FROM clause evaluation
  - 1. FROM clause: compute *cross-product* of Students and Enrolled
  - 2. WHERE clause: Check conditions, discard tuples that fail
  - 3. SELECT clause: Delete unwanted fields